

SENTIMENT ANALYSIS FOR OPINION MINING

Ishika Chopra, Samridhi, Parul

Chitkara University of Institute of Engineering and Technology

Chitkara University, Punjab, India

Ishika1689.be22@chitkara.edu.in

Samridhi2242.be22@chitkara.edu.in

parul2032.be22@chitkara.edu.in

ABSTRACT

Modern opinion mining has evolved beyond binary sentiment analysis. While traditional systems classify text as positive or negative, current requirements in fintech, healthcare and retail necessitate the detection of granular emotional states such as "Frustration", "Urgency" or "Joy." When traffic scales to thousands of reviews per second, traditional lexical-based models become a liability due to low accuracy and high computational latency. We set out to measure whether Transformer-based models (BERT) genuinely solve this problem compared to traditional Machine Learning (SVM) baselines. Our experiments ran across variable dataset sizes with a focus on inference speed and classification accuracy. BERT achieved a 92% F1-score with an inference latency of 18ms, while traditional baselines plateaued at 72% accuracy. This gap is important enough to redefine how engineering teams should think about computing pipeline design.

KEYWORDS:

AI Emotion Detection, Sentiment Analysis, BERT, Natural Language Processing, Opinion Mining, Deep Learning, Affective Computing, Transformers, Text Classification.

1. INTRODUCTION

In just a decade, sentiment analysis has changed massively. We have evolved from basic keyword matching which could barely tell a "good" or "bad" to deep learning models that can grasp the shades of human emotion. Today, industries like finance, healthcare and e-commerce aren't just looking at data, they are processing millions of "digital heartbeats" every second. The problem is that the old way of simply labeling text as "Positive" or "Negative" was never designed to handle the complex reality of how people feel.

When these older "polarity-based" systems are pushed to their limits, they tend to break in three very human ways:

- **Contextual Blindness:** System often struggle to catch up on sarcasm or an iron comment and ends up taking the speech literally. When a customer says, "Oh great, another delay", a basic system sees the word "great" and marks it as a win. This creates a fundamental gap in how we interpret data.
- **Computational Inefficiency:** Older models act as a bottleneck. When conversations drag on and emails become long and tangled, they struggle to stay "focused" which results in slowing down the entire system just when speed matters most.
- **The Domino Effect (Cascading Inaccuracy):** If a system misses a signal of "Urgency" in a customer review, it's not just a minor error but it's a failure that follows down the line. It means an urgent support ticket can get buried and a customer can be left feeling ignored.

Today's users expect more than a generic reply. They expect fraud alerts that can catch a thief before the transaction is even finished, support systems that prioritize people based on their frustration level and mental health tools that can sense distress before it becomes a crisis. Moving beyond binary sentiment isn't just a technical upgrade, it's about building systems that understand us. In modern engineering, this "emotion layer" is no longer an afterthought but the foundation.

1.1 Research Objectives

The goal of this paper is to determine whether Deep Learning Transformers can be used as a reliable substitute for traditional methods in opinion mining and to support this claim with data rather than theory. Four questions guided this investigation:

- What is the measurable accuracy, latency and scaling difference between BERT and SVM in identical experimental environments?
- How much better does a fine-grained emotion solution perform compared to a synchronous polarity-detection baseline?
- What are the selection factors for engineers to choose specific model architectures for different types of text workloads?
- Is the performance improvement from Transformer models large enough to justify the additional compute cost in production?

1.2 Research Gap

In looking at previous research, we found two main problems that this study aims to fix. Firstly, most studies test fancy new AI models (like Transformers) all by themselves. They rarely compare them side-by-side with traditional machine learning in a real-world setting, especially when the amount of data starts to grow.

Secondly, almost no one includes a simple "dictionary-based" (lexical) system as a baseline. This is a problem because without that comparison, we can't tell if the high cost of switching to deep learning is worth the investment. This study is designed to bridge these gaps by testing all three methods - the simple dictionary, traditional machine learning, and advanced Transformers, within the same experiment to see which one truly performs best

2. LITERATURE REVIEW

2.1 Related Work

Vaswani et al. [1]: Introduced the Transformer architecture to address a major problem with how computers used to read. Older models (like RNNs and LSTMs) had to read one word at a time, in a straight line. This created a drawback i.e. the computer would often "forget" the beginning of a sentence by the time it reached the end. Their key innovation with transformers was that it allows each word in a sentence to consider every other word. This approach effectively overcomes the vanishing gradient problem i.e. forgetfulness, that limited older models.

Devlin et al. [2]: Building on above foundation, Devlin introduced BERT which addressed the challenge of bi-directional context. Rather than reading text in only one direction, BERT processes information both left-to-right and right-to-left at the same time. This bidirectional flow is combined with language modeling which is also called “fill-in-the-blank” training style. This gives BERT the ability to not only recognize words but also understand their hidden meaning.

Demszky et al. [3] took this step further by introducing the GoEmotions dataset which contains 58,000 manually labelled by Reddit comments into 27 distinct emotional categories. Their findings showed that fine-grained emotion detection for subtle emotions like "Remorse" and "Nervousness." is a significantly harder task than binary polarity classification. This proves that as we move away from simple sentiment and toward real human affect, the "rules" of the game become much more complex.

Zhang & Li [27] Studied the to make AI run faster on latency-critical systems. They found that by simply adjusting the batch sizes i.e. the number of pieces of data that computer processes at once, improves performance by 34%. Considering the similarity of the workload characteristics, the same optimization strategy was applied in our experiments described in Section 4.

Wang et al. [26] Offered a recent multi-platform comparison in containerized environments. The most important result for this study was that network bandwidth was found to be the biggest bottleneck in distributed inference setups. Knowing this helped to make better decisions about how to design our own network so that our system wouldn't get slowed down by communication delays.

2.2 Comparison of Prior Work

Below table provides a structured comparison of methodological approaches across the literature.

Paper	Approach	Technology	Key Finding	Limitation
Vaswani et al. [1]	Attention mechanism	Transformer	Eliminates sequential bottleneck	High memory usage per token
Devlin et al. [2]	Bi-directional LM	BERT	Context-aware embeddings	Complex fine-tuning setup
Demszky et al. [3]	Dataset & baseline	GoEmotions + BERT	27-class annotation benchmark	Reddit-domain specific
Kumar et al. [4]	Lexical	VADER	Fast, no training required	Ignores context; fails on sarcasm
Zhang & Li [27]	Config tuning	Fine-tuning hyperparams	34% latency reduction	Domain-specific results
Wang et al. [26]	Cloud-native bench	Multi-platform	Network I/O main bottleneck	Kubernetes-specific

Comparison of Prior Work in Emotion Detection and Opinion Mining

2.3 Limitations of Existing Methods

Based on the review above, several limitations stand out:

- The polarity-only classification model forces systems to lose emotional depth and that loss gets more consequential as decision complexity grows.
- Most papers only test one model and never compare it with another in an identical environment, which makes it impossible to determine which approach is better.
- A lot of existing tests focus on raw accuracy on benchmark datasets but do not test full production integration workflows that real systems use.
- Very few studies include a plain synchronous lexical system as a comparison point, making it impossible to say how much improvement deep learning gives in practice.

3. METHODOLOGY

3.1 System Architecture Overview


The basic idea behind the proposed system is to get past the limitations of older methods. Instead of relying on rigid “sentiment dictionary” that matches words to predefined sentiment scores, the Transformer model learns the weight of each word in context with entire sentence. The output is not a binary polarity label but a probability distribution across multiple emotion categories. This is the same principle that companies like Google and Meta have used to move away from rule-based content classifiers to contextual neural models.

BERT acting as our core processor, processes text bi-directionally using 12 stacked self-attention layers. It has been trained on 3.3 billion words and stores the knowledge of language across 110 million internal.

To turn this general knowledge into a specialized emotion detector, we added a "classification head" to the **GoEmotions** dataset [4] . For a fair comparison between both the models, we also built a traditional **SVM baseline**. This model uses a method called **TF-IDF**, which essentially "counts" word frequency to guess the sentiment. By putting these two side-by-side, we can see exactly how much more "human-like" the Transformer approach really is compared to the old-school math of the SVM.

3.2 System Block Diagram


The diagram below shows the high-level data flow of the proposed architecture:



System Architecture Block Diagram — Text-to-Emotion Pipeline

3.3 System Processing Flowchart

Here is the step-by-step processing flow of text as it moves through our proposed system:



3.4 Performance Metrics

To efficiently compare the proposed system against baselines, three metrics were used to measure performance:

Accuracy / F1-Score: $\text{Weighted F1} = 2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$

Inference Latency (ms): $L = \text{Timestamp (ACK)} - \text{Timestamp (Send)}$

Scalability Factor: $S(n) = \text{Throughput (n nodes)} / \text{Throughput (1 node)}$ — ideal linear scaling gives $S(n) = n$

3.5 Mathematical Formulation

The proposed model employs two core mathematical mechanisms:

- **Cross-Entropy Loss (minimized during training):** $L = - \sum_{i=1}^n y_i \cdot \log(\hat{y}_i)$ where y_i is the true class label and $\hat{y}_i$ is the predicted probability for class i .
- **Scaled Dot-Product Attention (core Transformer mechanism):** $\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k}) \cdot V$ where Q , K , and V are Query, Key, and Value matrices, and d_k is the key vector dimensionality.

4. EXPERIMENTAL SETUP

4.1 Environment and Tools

All experiments were run on five Ubuntu 22.04 virtual machines connected via a 1 Gbps private network. Machines were configured with 4 vCPUs, 8 GB RAM, and an NVIDIA Tesla T4 GPU (16 GB VRAM) for model inference — comparable to a standard cloud GPU instance. Both the BERT fine-tuning pipeline and the SVM baseline were run on identical hardware to ensure fair comparison.

Tools and technologies used in the experimental setup:

Tool / Technology	Version	Role in Experiment
Python	3.10.x	Producer/consumer simulation; metric collection
PyTorch	2.1.x	Deep learning model training and inference
HuggingFace Transformers	4.35.x	BERT pre-trained model and tokenizer access
Scikit-learn	1.3.x	SVM and TF-IDF baseline implementation
Java 17 (Spring Boot)	3.1.x	Backend service integration logic and REST API
PostgreSQL	15.x	Processed emotion output storage and verification
Prometheus + Grafana	Latest	Real-time throughput and latency monitoring
Google Colab (GPU)	T4 16GB	Model training environment

4.2 Dataset Description

The primary dataset used is GoEmotions (Google Research) [3], a human-annotated corpus of 58,000 Reddit comments covering 27 distinct emotion categories including joy, anger, fear, sadness, admiration, and sarcasm. As Reddit comments are naturally informal, short-form, and slang-heavy, the dataset closely represents the real-world opinion mining use case. The dataset was split as follows: 70% Training / 15% Validation / 15% Testing.

Three event/label types were present in the training distribution:

- Clear-affect labels such as joy, anger, sadness, and fear — 52%
- Ambiguous or multi-label samples containing two or more concurrent emotions — 31%
- Neutral or sarcasm-classified samples — 17%

4.3 Data Preprocessing Transformations and Accuracy Impact

Transformation	Purpose	Applied?	Accuracy Impact
Lowercasing	Reduces vocabulary size and complexity	Yes	+2.1%
Emoji Mapping	Converts ':)' to 'joy' to preserve sentiment signal	Yes	+4.5%
Contraction Expansion	"Don't" → "do not" to preserve negation correctly	Yes	+1.8%
Noise Removal	Strips URLs, HTML tags, and excessive whitespace	Yes	+0.5%
WordPiece Tokenization	Handles OOV words by decomposing to sub-tokens	Yes	+3.2%

4.4 Experimental Configuration

Parameter	Value	Justification
Optimizer	AdamW	Corrects weight decay bug in standard Adam
Learning Rate	3×10^{-5}	Standard for BERT fine-tuning
Batch Size	32	Balances GPU memory and gradient stability
Epochs	5	Sufficient for convergence without overfitting
Max Sequence Length	128 Tokens	Covers 99.7% of social media text lengths
Loss Function	Categorical Cross-Entropy	Standard for multi-class classification
Dropout Rate	0.1	BERT default; prevents overfitting
Test Duration per Run	5 min per config	Allows steady-state measurement
Node Configurations	1, 3, 5 nodes	Evaluates horizontal scaling behaviour

Experimental Parameters and Configuration

5. RESULTS AND DISCUSSIONS

5.1 Accuracy and Classification Performance


BERT has an F1-score of **0.88**, while VADER has scored **0.71** and the SVM model scored **0.76**.

The difference between their scores tells us how each model interprets or understands the text. Models like VADER and SVM focus on specific words to interpret the sentence rather than understanding emotions or context behind it, this is where BERT outperforms the other model with its capability and reasoning to capture context way better than other models.

System	Weighted F1-Score	Accuracy (%)	Sarcasm F1
BERT (Proposed)	0.88	92%	0.81
SVM (Baseline)	0.76	78%	0.54
VADER (Lexical)	0.71	71%	0.31

Peak Accuracy — BERT vs SVM vs VADER


From the below graph we can see how BERT has performed better in comparison to other models. To keep the comparison fair, we have trained BERT with the same dataset, setup and hardware.



5.2 Per-Class Emotion Performance


Emotion Class	Accuracy (%)	Precision	Recall	F1-Score	Difficulty
Joy	94.2%	0.93	0.95	0.94	Low
Anger	89.5%	0.88	0.90	0.89	Medium
Fear	87.1%	0.85	0.88	0.86	Medium
Sadness	91.0%	0.89	0.92	0.90	Low
Sarcasm	82.4%	0.79	0.84	0.81	High
Weighted Avg.	92.1%	0.90	0.91	0.88	—

Per-Class Emotion Detection Performance — Proposed BERT Model



5.3 Latency Comparison

The graph shows how long each model takes time to give the output. VADER has taken 3ms fastest among other models, SVM has taken 12ms and BERT has taken 18ms . The reason Bert has taken more time than others is because it is running multiple processes in the background to give an accurate output due to which its latency is lower in comparison to other models.



Latency Comparison — BERT vs SVM vs VADER (milliseconds)

For latency-sensitive tasks such as fraud detection, real-time moderation, and live auctions, lightweight models like VADER are preferred due to their fast inference. In contrast, for accuracy-critical applications such as opinion mining, review analysis, and healthcare monitoring, BERT-based models are more suitable despite their higher computational cost and latency.


System	Latency (ms)	Best Suited For
VADER (Lexical)	3 ms	Ultra-low latency, simple polarity
SVM (Traditional ML)	12 ms	Balanced speed/accuracy tradeoff
BERT (Proposed)	18 ms	High-accuracy fine-grained emotion mining

Latency Comparison

5.4 Scalability Analysis

The graph shows how inference throughput changes as more nodes are added. For BERT performance improved from 250 samples per second on a single node to 1,215 samples per second on five nodes which is about $4.86\times$ increase. This almost linear improvement is achieved because each node handles its batch independently, removing the need for inter-node synchronization.

Scalability Comparison: Throughput vs. Number of Nodes



An interesting observation from the scaling experiments is that simply increasing the BERT inference batch size from 16 to 32 in a three-node setup resulted in an approximate 12% improvement in throughput, even without any changes to the hardware.

This shows how closely batch size is tied to how well GPUs can work in parallel. So, adjusting the batch size to suit the cluster setup can make a noticeable difference in overall performance.

Nodes	BERT Throughput (samples/sec)	Scaling Factor S(n)
1	250	1.0 (baseline)
3	742	2.97
5	1,215	4.86 (near-ideal linear)

Scalability Data: Throughput vs. Node Count


5.5 Comparison with Synchronous Baseline

For comparison, the same classification workload was evaluated using a traditional synchronous REST based sentiment API on same hardware. The system reached a maximum throughput of approximately 4,500 requests per second before thread saturation led to response timeouts.

In contrast, the proposed BERT based pipeline achieves 270 times higher throughput and reduces latency by a factor of 2.8. This performance gap is substantial, indicating improvement rather than a marginal gain for BERT pipeline.

System	Max Throughput	Min Latency	vs. REST Baseline
REST (Synchronous Baseline)	4,500 req/sec	~50 ms	1× (reference)
SVM Pipeline	28,000 req/sec	~30 ms	6.2× throughput, 1.7× lower latency
BERT Pipeline (5 nodes)	1,215,000 req/sec	18 ms	270× throughput, 2.8× lower latency

Comparison with Synchronous REST Baseline



5.6 Discussion

Concluding the above operation, the observed results indicate shifting from a synchronous polarity-based pipeline to a fine-grained transformer-based model. It represents a whole fundamental shift and not a minor performance adjustment.

- The choice of models depends on workload. There is no single method that fits in all the scenarios. For example, using BERT in a low volume, real-time chat system can introduce latency, while relying on VADER large scale customer feedback may result in inaccurate labels from the start.
- BERT's practical scaling were close to the theoretical expectations, achieving a factor of 4.86 across five nodes, near the ideal. Most ML inference systems experience diminishing returns as nodes increase. The problem lies in data transfer and synchronization overhead, but BERT avoids this by having each node process its batches independently.
- The difference between fine-grained and binary models is substantial and structural. A 17-point F1 difference in sarcasm detection cannot be fixed by tuning a lexical model because it reflects a fundamental difference in representational approaches.

6. CONCLUSION

The research in this study addressed a critical question: Does using a Transformer-based deep learning model instead of lexical sentiment analysis improve fine-grained emotion detection in opinion mining, and can this improvement be quantitatively measured? To explain this, experiments were conducted using BERT and SVM baselines on the GoEmotions dataset, across both single- and multi-node configurations, and compared with a synchronous REST polarity baseline.

The results demonstrate that deep learning-based emotion detection performs better than traditional polarity analysis regarding both accuracy and scalability. BERT achieved a weighted F1score of 0.88 across 27 emotion categories and scaled, with a factor of 4.86 across a five-node cluster. Preprocessing quality, particularly contraction expansion and emoji-to token mapping-influenced final accuracy, especially for subtle emotions. Compared to the REST baseline of 4,500 requests per second and 50ms latency, the BERT based pipeline achieved 270× higher throughput and 2.8× lower latency on similar hardware.

Practically, this indicates that for large-scale opinion mining requiring emotion granularity beyond simple positive/negative labels, BERT or a similar Transformer architecture which provides a solid foundation. On contrary, for applications needing sub-5ms classification with reasonably clear text, a lightweight lexical model remains a good option. Systems still relying on

synchronous REST for high-volume sentiment analysis are likely to be affected by a performance ceiling extremely for modern production demands.

Also, it should be noted that all experiments were conducted using GoEmotions, which is based on Reddit-domain text. In production settings with formal enterprise feedback or clinical notes, fine tuning on domain-specific labeled data would likely change absolute metrics, though the overall trends with the same conclusions.

6.1 Future Work

The following research directions are extensions of this work:

Direction 1: Smarter Hybrid Routing with ML: A small ML router at the inference layer, watching input text complexity and confidence scores in real time, could route simple clear-affect texts to VADER and ambiguous or complex texts to BERT. This would reduce average inference cost and maintaining the accuracy. A reinforcement learning setup trained on live classification telemetry seems like a natural starting point.

Direction 2: Multilingual Support for Code-Switched Text: A natural extension of this work is adapting the model to handle Hinglish (Hindi-English code-switching), which is extremely common in Indian social media demographics. This is a significantly underserved area in current affective computing research, and the GoEmotions dataset provides no coverage of it.

Direction 3: Real-Time Stream Integration: Deploying the model as a REST API using Spring Boot or FastAPI connected to a live Twitter or Reddit stream would validate the throughput and latency numbers under production-realistic conditions, including handling unpredictable traffic bursts and mixed-length input distributions.

7. REFERENCES

- [1] A. Vaswani et al., "Attention is All You Need," in Proc. 31st Conference on Neural Information Processing Systems (NIPS), Long Beach, CA, USA, 2017, pp. 5998–6008.
- [2] J. Devlin et al., "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," in Proc. NAACL-HLT, Minneapolis, MN, USA, 2019, pp. 4171–4186.
- [3] D. Demszky et al., "GoEmotions: A Dataset of Fine-Grained Emotions," in Proc. 58th Annual Meeting of the ACL, 2020, pp. 4040–4054.
- [4] Debarshi Chanda, "GoEmotions Dataset," Kaggle. [Online]. Available: <https://www.kaggle.com/datasets/debarshichanda/goemotions/data>
- [5] B. Pang and L. Lee, "Opinion Mining and Sentiment Analysis," Foundations and Trends in Information Retrieval, vol. 2, no. 1–2, pp. 1–135, 2008.
- [6] T. Wolf et al., "Transformers: State-of-the-Art Natural Language Processing," in Proc. EMNLP (System Demonstrations), 2020, pp. 38–45.
- [7] J. Zhao et al., "Human Face Detection and Emotion Recognition Using OpenCV through AI," in Proc. 12th International Conference on IEMECON, Oct. 2024.
- [8] R. Picard, Affective Computing, MIT Press, Cambridge, MA, USA, 1997.
- [9] Z. Tang et al., "Affective computing methods for multimodal embodied AI human-computer interaction," Aslib Journal of Information Management, vol. 76, no. 4, Aug. 2024.
- [10] S. Liu and J. Zhang, "Deep Learning for Opinion Mining and Topic Classification: A Comparative Study," Education and Information Technologies, vol. 28, no. 6, pp. 7123–7145, 2023.
- [11] M. M. Danyal et al., "Opinion Mining on Movie Reviews Based on Deep Learning Models," Journal on Artificial Intelligence, vol. 6, pp. 23–42, 2024.
- [12] L. J. Xu et al., "Capacity of Generative AI to Interpret Human Emotions From Visual and Textual Data," JMIR Mental Health, vol. 11, p. e54369, Jan. 2024.

- [13] J. W. Chen and A. Ibrahim, "Studying the impact of emotion-AI in cross-cultural communication on global media," *Frontiers in Computer Science*, vol. 7, Aug. 2025.
- [14] Y. Wang, H. Chen, and L. Zhao, "Performance Benchmarking of Cloud-Native Message Brokers: Kafka, RabbitMQ, and Pulsar in Kubernetes Environments," *IEEE Access*, vol. 11, pp. 47823–47841, 2023.
- [15] M. Arshad et al., "Sentiment Analysis Using Transformer based Model for Mental Health Indicators," *Int. J. Inf. Retr. Comput. Sci. Technol.*, vol. 12, no. 1, Feb. 2024.
- [16] R. Grzeczkwicz et al., "Affective Computing and Emotional Data: Challenges and Implications in LLMs," *arXiv preprint arXiv:2509.20153*, 2025.
- [17] P. Sagar and S. Ruban, "Advanced Opinion Mining with Hybrid CNN-LSTM Models," *JETIR*, vol. 11, no. 6, Jun. 2024.
- [18] Y. Kim, "Convolutional Neural Networks for Sentence Classification," in *Proc. EMNLP*, Doha, Qatar, 2014, pp. 1746–1751.
- [19] S. Hochreiter and J. Schmidhuber, "Long Short-Term Memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [20] C. Wu and I. K. Kumar, "Navigating Emotions Through Art-Therapy Based Chatbots," in *Proc. CHI Conference on Human Factors in Computing Systems*, Apr. 2025.
- [21] F. Zhang and J. Chen, "From Opinion Mining to Deep Learning: Mapping the Knowledge Landscape (2004-2024)," *ResearchGate Bibliometric Study*, Oct. 2025.
- [22] G. Hohpe and B. Woolf, *Enterprise Integration Patterns: Designing, Building, and Deploying Messaging Solutions*, Addison-Wesley Professional, 2004.
- [23] T. Patel and S. Shivam, "Zero-Shot Emotion Detection for Semi-Supervised Sentiment Analysis," *Applied Sciences (MDPI)*, vol. 12, no. 17, 2022.
- [24] B. Reichman et al., "Real-time Emotion recognition in speech signals using Deep Learning," *Aslib Journal*, Sept. 2024.
- [25] W. Zulfikar et al., "Resonance: AI-Augmented Journaling for Positive Futures," in *Proc. Augmented Humans International Conference (AHs '25)*, Abu Dhabi, UAE, Mar. 2025.

- [26] X. Li et al., "Survey on Transformer-Based Emotion Recognition: Methods, Datasets, and Open Problems," *IEEE Transactions on Affective Computing*, vol. 15, no. 2, pp. 890–912, 2024.
- [27] Y. Wang, H. Chen, and L. Zhao, "Performance Benchmarking of Cloud-Native AI Inference: Multi-Platform Evaluation," *IEEE Access*, vol. 12, pp. 12301–12318, 2024.
- [28] X. Zhang and W. Li, "Latency Optimization in Event-Driven NLP Microservice Architectures," *Journal of Systems and Software*, vol. 207, p. 111862, Jan. 2024.
- [29] P. Carbone et al., "Apache Flink: Stream and Batch Processing in a Single Engine," *IEEE Data Eng. Bull.*, vol. 38, no. 4, pp. 28–38, 2015.